

# An Introduction to Deep Learning

Clint P. George

University of Florida Informatics Institute


March 14, 2017

---

## Outline

- ① Introduction to feed-forward networks
- ② Relation to logistic regression
- ③ Notes on implementation
- ④ Illustration using synthetic and real data

# Motivation

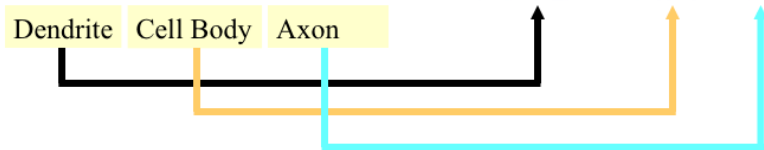
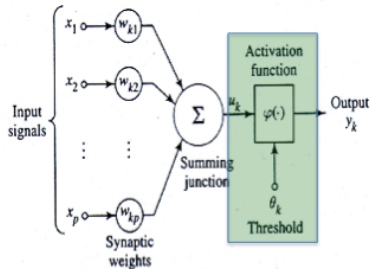
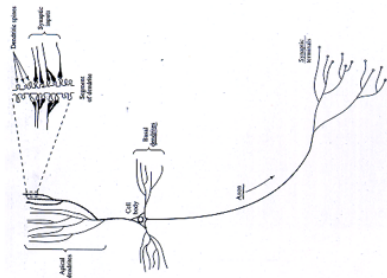


The idea: Human intelligence may be due to a learning algorithm.  
We aim to build algorithms that mimic the brain.<sup>1</sup>

---

<sup>1</sup>image: <https://backyardbrains.com/experiments/EEG>

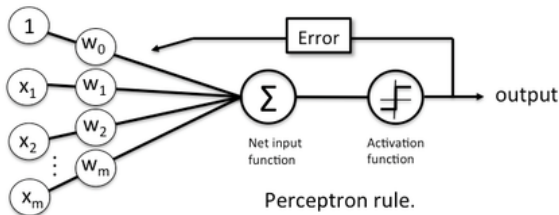
# Imitate neurons in the brain: Artificial Neurons



Artificial Neuron (AN): input, weights, and output

## Activation functions

The Perceptron (Rosenblatt et al. 1957 & 1962) computes a step function as an activation function.



$$\text{step}(z) = \begin{cases} 1 & z \geq t \\ 0 & z < t \end{cases}, \text{ where } t \text{ is a threshold}$$

As each input is applied to the perceptron its output is compared to the target. To keep the output closer to the target the **learning rule** adjusts the network parameters.

## What can a single AN compute?

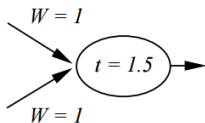
The Perceptron output is given by  $y = \text{step}(b + \sum_{j=1}^p w_j x_j)$ .

Perceptron can divide the input space into two regions.

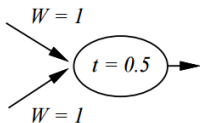
The decision boundary is given by:  $b + \sum_{j=1}^p w_j x_j = 0$ .

Perceptron can learn to classify any linearly separable set of inputs—convergence theorem (Rosenblatt 1962)

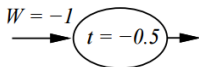
# Boolean functions



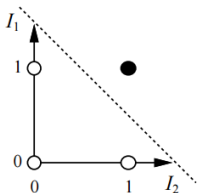
**AND**




**OR**



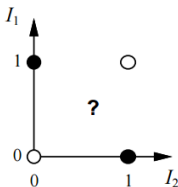
**NOT**



(a)  $I_1$  and  $I_2$



(b)  $I_1$  or  $I_2$



(c)  $I_1$  xor  $I_2$

Examples<sup>2</sup> of linearly separable and non separable problems

---

<sup>2</sup>Veloso, 2001

## Learning the XOR function

Consider this as a regression problem and use the MSE loss function:

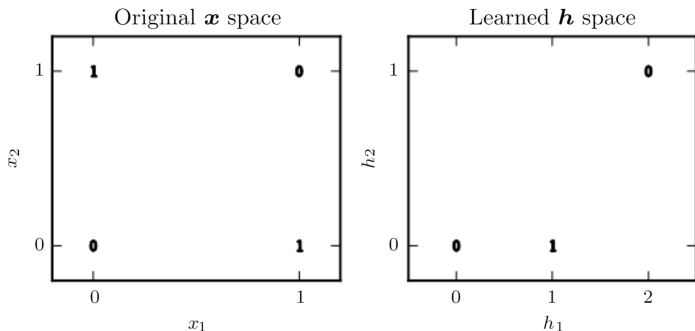
$$\text{MSE}(\theta) = \frac{1}{4} \sum_{\mathbf{x} \in \mathcal{X}} (y - f_{\theta}(\mathbf{x}))^2$$

where  $\theta = (\mathbf{w}, b)$  and  $f_{\theta}(\mathbf{x}) = b + \sum_{j=1}^p w_j x_j$ —a linear model.

Using the normal equations we can minimize  $\text{MSE}(\theta)$  w.r.t.  $\mathbf{w}$  and  $b$  in closed form. It gives  $\mathbf{w} = 0$  and  $b = .5$ —This gives the model output .5 everywhere.



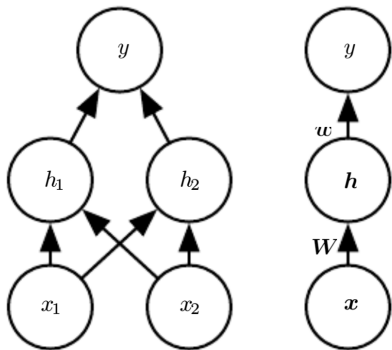
# Learning the XOR function



- When  $x_1 = 0$ , the model's output must increase as  $x_2$  increases
- When  $x_1 = 1$ , the model's output must decrease as  $x_2$  increases
- A linear model applies a fixed coefficient  $w_2$  to  $x_2$ . It cannot use the value of  $x_1$  to change the coefficient  $w_2$  on  $x_2$  and cannot solve this problem.

## Intuition behind multilayer neural network

One way to solve the XOR problem is to transform the input by introducing a feed-forward network:



The complete model will then be, in a function form:

$$f(\mathbf{x}; \mathbf{W}, \mathbf{w}, b, c) = f^{(2)} \left( f^{(1)}(\mathbf{x}; \mathbf{W}, c); \mathbf{w}, b \right)$$

What function should  $f^{(1)}$  be?

We consider  $f^{(2)}$  as a linear function.

We can write the hidden layer output as

$$\mathbf{h} = f^{(1)}(\mathbf{x}; \mathbf{W}, \mathbf{c}) = g(\mathbf{W}^\top \mathbf{x} + \mathbf{c})$$

## What function should $f^{(1)}$ be?

We consider  $f^{(2)}$  as a linear function.

We can write the hidden layer output as

$$\mathbf{h} = f^{(1)}(\mathbf{x}; \mathbf{W}, \mathbf{c}) = g(\mathbf{W}^T \mathbf{x} + \mathbf{c})$$

If  $f^{(1)}$  is also linear, then the network as a whole would remain a linear function of its input.

## What function should $f^{(1)}$ be?

We consider  $f^{(2)}$  as a linear function.

We can write the hidden layer output as

$$\mathbf{h} = f^{(1)}(\mathbf{x}; \mathbf{W}, \mathbf{c}) = g(\mathbf{W}^T \mathbf{x} + \mathbf{c})$$

So we must use a nonlinear activation function for  $g$ . A popular nonlinear activation function  $g$  is rectified linear unit (ReLU)

$$g(z) = \max\{0, z\}$$

## A feed-forward network solution to XOR

We wish the network to perform well on the four cases

$$\mathbf{X} = \begin{bmatrix} 0 & 0 \\ 0 & 1 \\ 1 & 0 \\ 1 & 1 \end{bmatrix}, \mathbf{y} = \begin{bmatrix} 0 \\ 1 \\ 1 \\ 0 \end{bmatrix}$$

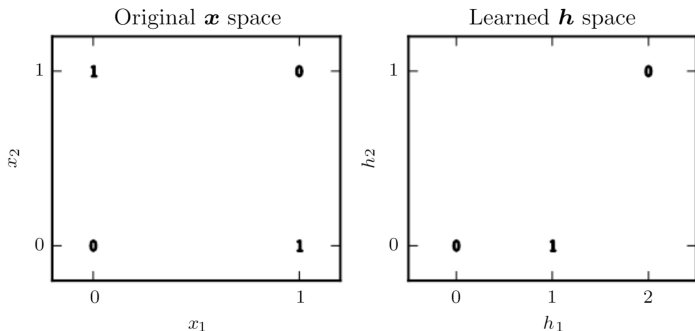
We can specify a two-layer network solution to XOR as

$$\mathbf{W} = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}, \mathbf{c} = \begin{bmatrix} 0 \\ -1 \end{bmatrix}, \mathbf{w} = \begin{bmatrix} 1 \\ -2 \end{bmatrix}, b = 0$$

The complete network is given by

$$f(\mathbf{x}; \mathbf{W}, \mathbf{w}, b, \mathbf{c}) = \mathbf{w}^T \max\{0, \mathbf{W}^T \mathbf{x} + \mathbf{c}\} + b$$

## A feed-forward network solution to XOR



In the proposed network, the nonlinear hidden layer has mapped both  $\mathbf{x} = [1, 0]$  and  $\mathbf{x} = [0, 1]$  to a single point in feature space,  $\mathbf{h} = [1, 0]$ .

A linear model can now describe the function as increasing in  $h_1$  and decreasing in  $h_2$ .

## Activation functions

Modern ANs use a variety of activation functions that are smoother than the step function.

Linear - no input squashing

$$y = x$$

Logistic sigmoid - squash input into  $[0, 1]$

$$y = \text{sigmoid}(x) = \frac{1}{1 + \exp(-x)}$$

Hyperbolic tangent - squash input into  $[-1, 1]$

$$y = \tanh(x) = \frac{\exp(x) - \exp(-x)}{\exp(x) + \exp(-x)}$$

Rectified linear unit

$$y = \max\{0, x\}$$



## Relation to logistic regression

Let  $p(y = 1 \mid X = \mathbf{x}) = p(\mathbf{x}; \mathbf{w})$  be the conditional probability that a particular sample belongs to class 1 given its predictors  $\mathbf{x}$ . We write the logistic regression model as

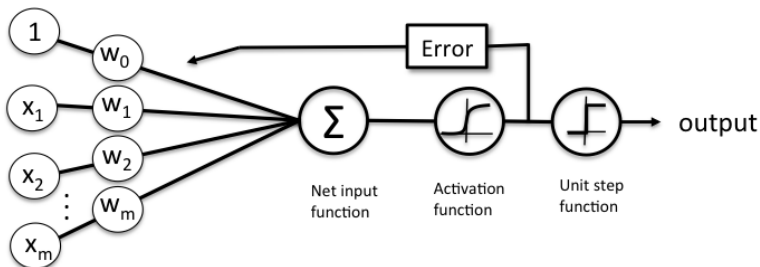
$$\text{logit}(p(\mathbf{x}; \mathbf{w})) = w_0 + \mathbf{w}^\top \mathbf{x}$$

Solving for  $p(\mathbf{x}; \mathbf{w})$  gives

$$p(\mathbf{x}; \mathbf{w}) = \text{sigmoid}(w_0 + \mathbf{w}^\top \mathbf{x})$$

To minimize misclassification rate, one should predict  $y = 1$  when  $p \geq .5$ , and vice versa.—i.e. guess 1 whenever  $w_0 + \mathbf{w}^\top \mathbf{x} \geq 0$  and 0 otherwise. So logistic regression gives a linear classifier.

## Relation to logistic regression



**Schematic of a logistic regression classifier.**

3

Learning parameters  $w$  and  $b$ : maximum likelihood estimation

- Perceptron algorithm: online and error-driven.
- Logistic regression: batch algorithms—e.g. gradient descent, limited-memory BFGS, or online algorithms—stochastic gradient descent.

---

<sup>3</sup>[http://rasbt.github.io/mlxtend/user\\_guide/classifier/LogisticRegression/](http://rasbt.github.io/mlxtend/user_guide/classifier/LogisticRegression/)

## Gradient-based learning

The nonlinearity of a neural network causes most interesting loss functions to become non-convex

- Optimization procedure – using iterative, gradient-based optimizers that drive the cost function to a very low value
- Cost function,  $C(\theta)$  – they are more or less the same as those for other parametric models, such as linear models

## Cost function

We define the loss functional  $\mathcal{L}(f_\theta, z)$  based on the network, e.g., squared error, the negative conditional log-likelihood. Here,  $z = (x, y)$  and  $f_\theta(x)$  is the predictive function for  $y$  given  $\theta$ .

We define the cost function as

$$C(\theta) = \int \mathcal{L}(f_\theta, z) P(z) dz$$

We typically write this as an average—*training loss*

$$C(\theta) = \frac{1}{n} \sum_{i=1}^n \mathcal{L}(f_\theta, z)$$

## Gradient descent (GD) algorithm<sup>4</sup>

We find a  $\theta$  that minimizes the cost

- By solving  $\frac{\partial C(\theta)}{\partial \theta} = 0$  we can find the minima, maxima, and saddle points.
- In general, we cannot find the solutions of this equation. So we seek numerical optimization methods
- *local descent*: iteratively modify  $\theta$  so as to decrease  $C(\theta)$ , until we reach a local minima

$$\theta^{(t+1)} = \theta^{(t)} - \epsilon \frac{\partial C(\theta^{(t)})}{\partial \theta^{(t)}},$$

$\epsilon$  is the learning rate

---

<sup>4</sup><http://www.iro.umontreal.ca/~pift6266/H10/notes/gradient.html>

## Stochastic gradient descent (SGD) algorithm

- We use the fact that  $C(\theta)$  is an average over i.i.d. samples
- Make updates much often

$$\theta^{(t+1)} = \theta^{(t)} - \epsilon \frac{\partial \mathcal{L}(\theta^{(t)}, z)}{\partial \theta^{(t)}},$$

$z$  is an the next sample from the training set.—It can be implemented online

- The update direction is a random variable whose expectation is the true gradient of interest.

## Output units

The output layer provides additional transformation from the hidden features  $\mathbf{h}$  to complete the network's indented task.

**Linear** units for Gaussian output distributions: the output units based on an affine transformation with no nonlinearity

- Given hidden features  $\mathbf{h}$ , we define outputs  $f_{\theta}(\mathbf{x}) = \mathbf{w}^{\top} \mathbf{h} + b$
- It's typically used to produce the mean of a conditional Gaussian distribution  $p(y|\mathbf{x}, \theta) = \mathcal{N}(f_{\theta}(\mathbf{x}), \mathbf{I})$

## Cost function: the negative log-likelihood

The neural network defines a conditional distribution  $p(y | \mathbf{x}, \theta)$ .

Suppose we use MLE for parameter estimation. Then, it's natural to use the cost function as the negative log-likelihood:

$$C(\theta) = -\mathbb{E}_{\mathbf{x}, y} \log p_{\text{model}}(y | \mathbf{x}, \theta)$$

### Example

if  $p_{\text{model}}(y | \mathbf{x}, \theta) = \mathcal{N}(f_{\theta}(\mathbf{x}), I)$ , then we have

$$C(\theta) = \frac{1}{2} \mathbb{E}_{\mathbf{x}, y} \|y - f_{\theta}(\mathbf{x})\|^2 + \text{const.}$$



## Output units

**Sigmoid** units for Bernoulli output distributions: e.g. a binary classification problem

Given hidden features  $\mathbf{h}$ , we define the output as

$$y = \text{sigmoid}(z) = \text{sigmoid}(\mathbf{w}^\top \mathbf{h} + b)$$

The sigmoid can be motivated by constructing an unnormalized probability distribution that doesn't sum to 1. We define a probability distribution over  $y$  using the value  $z$  as follows:

$$\log \tilde{p}(y | \mathbf{x}) = yz \tag{1}$$

$$p(y | \mathbf{x}) = \frac{\exp(yz)}{\sum_y \exp(yz)} \tag{2}$$

$$= \text{sigmoid}((2y - 1)z) \tag{3}$$

This yields a Bernoulli distribution controlled by a sigmoidal transformation of  $z$

## Parameter estimation via MLE

We assume we have a training set  $\mathcal{X} = (\mathbf{x}_i, y_i), i = 1, 2, \dots, n$ . Our goal is to maximize the log likelihood:

$$\log p(\mathcal{X} | \theta) = \sum_{i=1}^n \log p(\mathbf{x}_i, y_i | \theta).$$

We factor the log likelihood into an unconditional term  $p(\mathbf{x})$  that we ignore and a conditional term  $p(y | \mathbf{x})$  that we focus.

Let  $z_i = \mathbb{I}(y_i == 1)$  and  $p_i = p(y_i = 1 | \mathbf{x}_i)$ —a logistic linear function of  $\mathbf{x}_i$  for a single layer network. We can write the conditional density of the dataset as

$$\mathcal{L} = \prod_{i=1}^n p_i^{z_i} (1 - p_i)^{1-z_i}$$

## Parameter estimation via MLE



We wish to maximize the log likelihood:

$$\log \mathcal{L} = \sum_{i=1}^n z_i \log p_i + (1 - z_i) \log (1 - p_i)$$

The negative of this log likelihood is a *cross entropy* between the indicator variables  $z$  and the posterior probabilities  $p$ .

This also shows that the cross entropy is a natural cost function for binary classification problem (Jordan 1995)

## Hidden units: sigmoid and hyperbolic tangent<sup>5</sup>





Left: Common neural activation function motivated by biological data. Right: Logistic sigmoid and hyperbolic tangent

---

<sup>5</sup>Glorot (2011)

## Hidden units: sigmoid and hyperbolic tangent<sup>5</sup>





Sigmoidal units saturate across most of their domain—they saturate to a high value when  $z$  is very positive and vice versa. It can make gradient-based learning very difficult.

---

<sup>5</sup>Glorot (2011)

## Hidden units: sigmoid and hyperbolic tangent<sup>5</sup>



The hyperbolic tangent has a steady state at 0, hence is preferred from the optimization standpoint. It forces an antisymmetry around 0 which is absent in biological neurons.



---

<sup>5</sup>Glorot (2011)

## Hidden units: ReLU

A popular choice of activation function for hidden units<sup>6</sup>

$$g(z) = \max\{0, z\}$$





---

<sup>6</sup>Glorot (2011)

## Hidden units: ReLU

A popular choice of activation function for hidden units<sup>6</sup>

$$g(z) = \max\{0, z\}$$



*One can safely disregard nonlinearity (Goodfellow et al. 2016): “... neural network training algorithms do not usually arrive at a local minimum of the cost function, but instead merely reduce its value significantly”*

---



<sup>6</sup>Glorot (2011)



## Hidden units: ReLU

A popular choice of activation function for hidden units<sup>6</sup>

$$g(z) = \max\{0, z\}$$



Non-linearity in the network comes from the path selection associated with individual neurons being active or not. Once this subset of neurons is selected, the output is a linear function of the input.


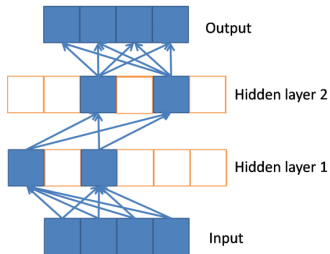
---

<sup>6</sup>Glorot (2011)

## Hidden units: ReLU

A popular choice of activation function for hidden units<sup>6</sup>

$$g(z) = \max\{0, z\}$$



ReLU allows a network to easily obtain sparse representations.

---

<sup>6</sup>Glorot (2011)

## An overview of back-propagation algorithm

Computing an analytical expression for the gradient is straightforward, but numerical evaluation of such an expression can be expensive.—**backprop** gives an inexpensive procedure to evaluate this gradient.

**backprop** uses the chain rule of calculus: Suppose  $\mathbf{x} \in \mathbb{R}^m$ ,  $\mathbf{y} \in \mathbb{R}^n$ ,  $g : \mathbb{R}^m \rightarrow \mathbb{R}^n$ ,  $f : \mathbb{R}^n \rightarrow \mathbb{R}$ . If  $\mathbf{y} = g(\mathbf{x})$ ,  $z = f(\mathbf{y})$ , then

$$\frac{\partial z}{\partial x_i} = \sum_j \frac{\partial z}{\partial y_j} \cdot \frac{\partial y_j}{\partial x_i}$$

We can apply this procedure to vectors and tensors in the multilayer network

## Notes on architectural considerations

In practice, the overall structure of the network is important: how many units it should have and how these units should be connected to each other.

Alternatives to feed-forward networks

- Convolutional Neural Networks — imitates human memory
- Auto Encoders — unsupervised learning and dimensionality reduction

# Hands on experiments

## Datasets

- A synthetic spiral dataset with multiple classes
- The MNIST data set with 0-9 handwritten characters

Algorithm: A basic back-propagation algorithm implementation with one hidden layer

Link to R scripts: <http://bit.ly/deep-learning-stats>

## Links to serious implementations

Deep learning frameworks, which uses CPU and GPU

- TensorFlow with R  
<https://rstudio.github.io/tensorflow/index.html>
- Theano with python  
<http://deeplearning.net/software/theano>

Available R packages

- neuralnet
- deepnet
- h2o

**Questions?**